

Fashion

May 4, 2025

```
[3]: # 1. Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
↳Dropout
from sklearn.metrics import classification_report, confusion_matrix
```

```
[4]: # 2. Load dataset from CSV
train_df = pd.read_csv("fashion-mnist_train.csv")
test_df = pd.read_csv("fashion-mnist_test.csv")

# 3. Split features and labels
X_train = train_df.iloc[:, 1:].values.reshape(-1, 28, 28, 1) / 255.0
y_train = train_df.iloc[:, 0].values

X_test = test_df.iloc[:, 1:].values.reshape(-1, 28, 28, 1) / 255.0
y_test = test_df.iloc[:, 0].values
```

```
[5]: # 4. Define CNN model
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D(2,2),
    Flatten(),
    Dropout(0.3),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax') # 10 categories
])

# 5. Compile model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
# 6. Train model
```

```
history = model.fit(X_train, y_train, epochs=10, validation_split=0.1,  
                    verbose=2)
```

c:\Users\darsh\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/10

1688/1688 - 19s - 11ms/step - accuracy: 0.8201 - loss: 0.4915 - val_accuracy: 0.8717 - val_loss: 0.3542

Epoch 2/10

1688/1688 - 17s - 10ms/step - accuracy: 0.8765 - loss: 0.3364 - val_accuracy: 0.8938 - val_loss: 0.3010

Epoch 3/10

1688/1688 - 16s - 9ms/step - accuracy: 0.8908 - loss: 0.2925 - val_accuracy: 0.8958 - val_loss: 0.2777

Epoch 4/10

1688/1688 - 16s - 10ms/step - accuracy: 0.9029 - loss: 0.2631 - val_accuracy: 0.8873 - val_loss: 0.3045

Epoch 5/10

1688/1688 - 16s - 10ms/step - accuracy: 0.9108 - loss: 0.2375 - val_accuracy: 0.9033 - val_loss: 0.2600

Epoch 6/10

1688/1688 - 16s - 10ms/step - accuracy: 0.9166 - loss: 0.2221 - val_accuracy: 0.9072 - val_loss: 0.2526

Epoch 7/10

1688/1688 - 16s - 9ms/step - accuracy: 0.9221 - loss: 0.2078 - val_accuracy: 0.9130 - val_loss: 0.2435

Epoch 8/10

1688/1688 - 17s - 10ms/step - accuracy: 0.9270 - loss: 0.1922 - val_accuracy: 0.9160 - val_loss: 0.2349

Epoch 9/10

1688/1688 - 16s - 10ms/step - accuracy: 0.9312 - loss: 0.1803 - val_accuracy: 0.9132 - val_loss: 0.2402


Epoch 10/10

1688/1688 - 16s - 10ms/step - accuracy: 0.9346 - loss: 0.1715 - val_accuracy: 0.9113 - val_loss: 0.2466

[6]: # 7. Plot training and validation accuracy

```
plt.plot(history.history['accuracy'], label='Train Accuracy')  
plt.plot(history.history['val_accuracy'], label='Val Accuracy')  
plt.title("Training vs Validation Accuracy")  
plt.xlabel("Epochs")
```

```
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



```
[7]: # 8. Evaluate on test set
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=0)
print(f"Test Accuracy: {test_acc:.2f}")
```

Test Accuracy: 0.92


```
[8]: # 9. Predict and evaluate
y_pred_probs = model.predict(X_test)
y_pred = np.argmax(y_pred_probs, axis=1)

# Confusion Matrix
conf_mat = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(10,8))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
plt.show()

# Classification Report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

313/313 1s 4ms/step



Classification Report:

	precision	recall	f1-score	support
0	0.90	0.81	0.85	1000
1	0.99	0.99	0.99	1000
2	0.90	0.82	0.86	1000
3	0.91	0.95	0.93	1000
4	0.86	0.89	0.87	1000

5	0.99	0.98	0.99	1000
6	0.74	0.81	0.77	1000
7	0.95	0.97	0.96	1000
8	0.99	0.99	0.99	1000
9	0.97	0.96	0.97	1000
accuracy			0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

```
[9]: # 10. Show 10 predictions vs actual labels
class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
               'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

plt.figure(figsize=(15, 5))
for i in range(10):
    plt.subplot(2, 5, i+1)
    plt.imshow(X_test[i].reshape(28, 28), cmap='gray')
    plt.axis('off')
    plt.title(f"True: {class_names[y_test[i]]}\nPred: {class_names[y_pred[i]]}")
plt.tight_layout()
plt.show()
```

